

Informed search

Classes of Search

Blind (uninformed)	Depth-First	Systematic exploration of whole tree until the goal is found.
	Breadth-First	
	Iterative-Deepening	
Heuristic (informed)	Hill-Climbing	Uses heuristic measure of goodness of a node, e.g. estimated distance to goal.
	Best-First	
	Beam	
Optimal (informed)	Branch&Bound	Uses path "length" measure. Finds "shortest" path. A* also uses heuristic
	A*	

INFORMED SEARCH STRATEGIES


- One that uses problem specific **knowledge beyond the definition of the problem** itself.
- Find solutions more efficiently than an uninformed strategy.

Best-first search

- Uses an **evaluation function** $f(n)$ for each node
 - $f(n)$ provides an estimate for the total cost.
 - Expand the node n with smallest $f(n)$.
- Implementation:**
Order the nodes in increasing order of cost.
- Special cases:
 - Greedy best-first search
 - A* search


- Most best-first algorithms include as a component of **f** a **heuristic function denoted by $h(n)$**
- $h(n)$** = estimated cost of the cheapest path from the **state at node n** to a **goal state**

A simplified road map of part of Romania



Romania with straight-line distance

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

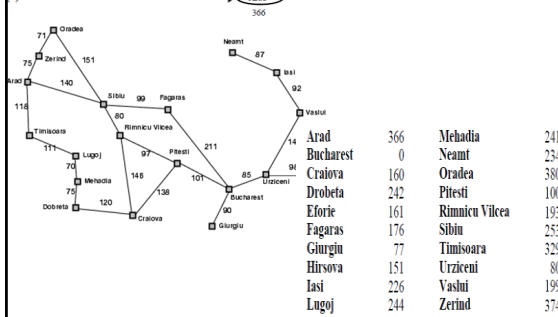
Values of h_{SLD} —straight-line distances to Bucharest

Greedy Best first search

- $f(n)$ = estimate of cost from n to goal
- $f(n)$ = straight-line distance from n to **Bucharest**
- Greedy best-first search expands the node that **appears to be closest to goal**.

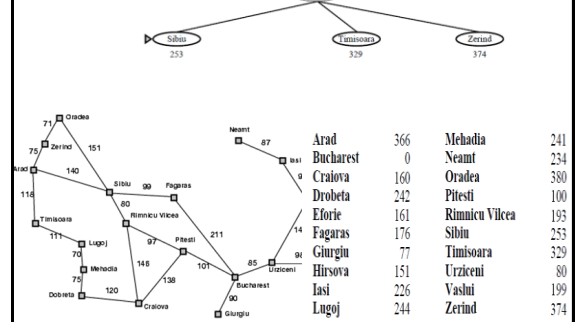
Greedy best first search example

a) The initial state

Values of h_{SLD} —straight-line distances to Bucharest.

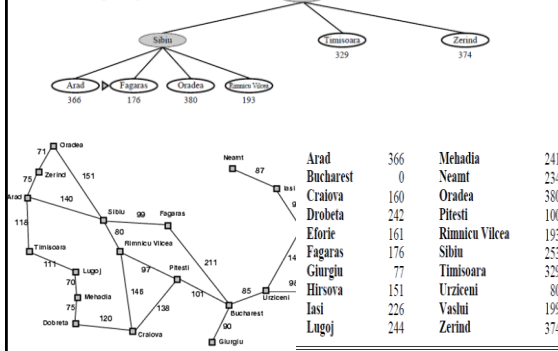
Greedy best first search example

(b) After expanding Arad



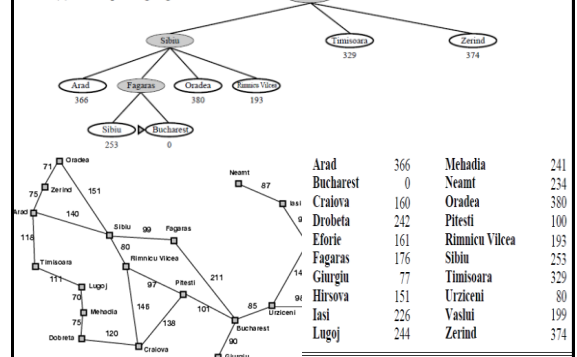
Greedy best first search example

(c) After expanding Sibiu



Greedy best first search example


(d) After expanding Fagaras



Properties of greedy best-first search

- Can stuck in loops.
- Optimal? No

e.g. Arad → Sibiu → Rimnicu Virea → Pitesti → Bucharest is shorter



A* search


- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
- $g(n)$ = cost so far to reach n
- $h(n)$ = estimated cost from n to goal
- $f(n)$ = estimated total cost of path through n to goal
- Best First search has $f(n)=h(n)$

Admissible heuristics

- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{\text{SLD}}(n)$ (never overestimates the actual road distance)

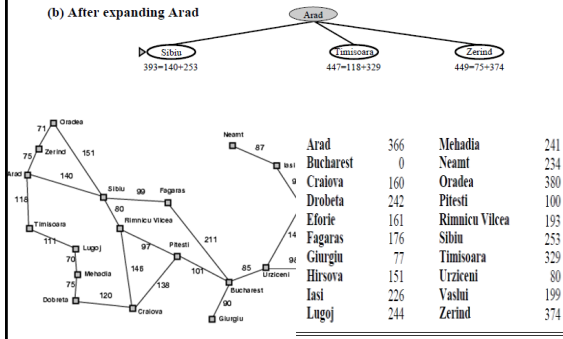
A* search example

(a) The initial state



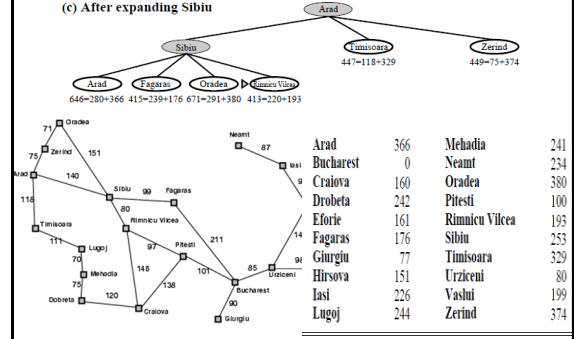
A* search example

(b) After expanding Arad



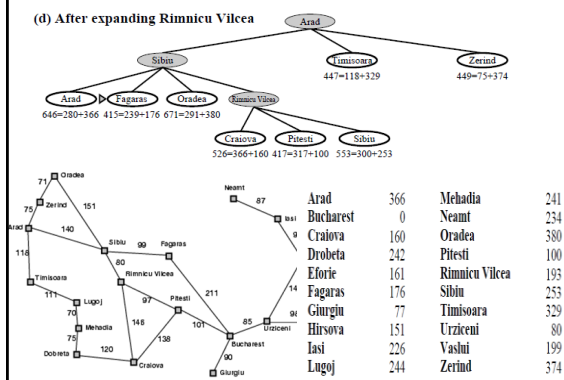
A* search example

(c) After expanding Sibiu



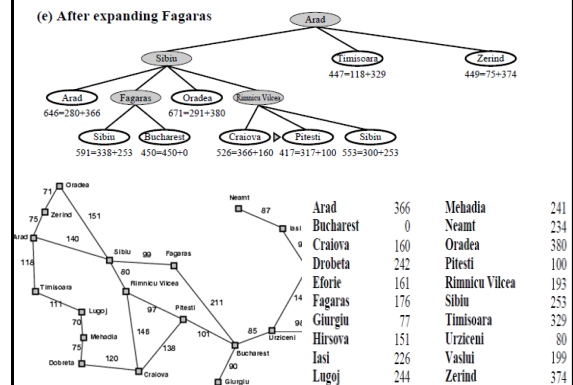
A* search example

(d) After expanding Rimnicu Vilcea



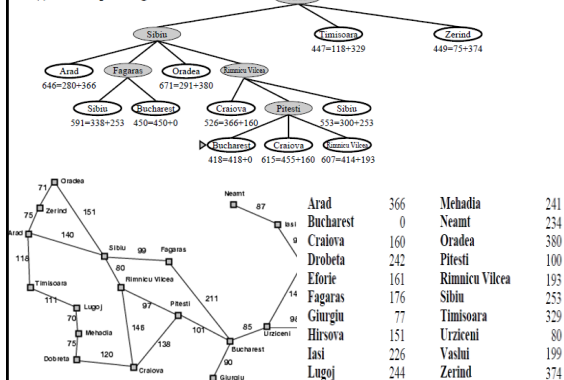
A* search example

(e) After expanding Fagaras



A* search example

(f) After expanding Pitesti




Hill-climbing search

- If there exists a **successor state s** for the **current state n** such that
 - $h(s) < h(n)$
 - $h(s) \leq h(t)$ for all the successors t of n ,
 then **move from n to s**. Otherwise, halt at n .
- Looks one step ahead to determine if any successor is better than the current state; if there is, move to the best successor.
- Similar to Greedy search, it uses h but **does not allow backtracking or jumping to an alternative path since it doesn't "remember"** where it has been.
- **Not complete** since the search will terminate at "local minima", "plateaus", and "ridges".

Hill-climbing search

- Problem: depending on initial state, can get stuck in local maxima



Hill-climbing search

- "Like climbing Everest in thick fog with amnesia (memory loss)"

function HILL-CLIMBING(*problem*) returns a state that is a local maximum
 inputs: *problem*, a problem

local variables: *current*, a node
neighbor, a node


current ← MAKE-NODE(INITIAL-STATE[*problem*])

loop do

neighbor ← a highest-valued successor of *current*

if VALUE[*neighbor*] ≤ VALUE[*current*] then return STATE[*current*]

current ← *neighbor*



Drawbacks of hill climbing

- Problems:
 - **Local Maxima:** peaks that aren't the highest point in the space
 - **Plateaus:** the space has a broad flat region that gives the search algorithm no direction
 - **Ridges:** flat like a plateau, but with dropoffs to the sides; steps to the North, East, South and West may go down, but a step to the NW may go up.
- Remedy:
 - Random restart.
- Some problem spaces are great for hill climbing and others are terrible.