

Red Black Tree

Red-black trees: Overview

- ♦ Red-black trees are a variation of binary search trees to ensure that the tree is **balanced**.
 - » Height is $O(\lg n)$, where n is the number of nodes.
- ♦ Operations take $O(\lg n)$ time in the **worst case**.

Applications

Red Black trees are used in

- ♦ Completely Fair Scheduler in Linux Kernel: completely fair scheduler, linux/rbtree.h
- ♦ Computational Geometry Data structures
- ♦ Java: java.util.TreeMap , java.util.TreeSet .
- ♦ C++ STL: map, multimap, multiset.

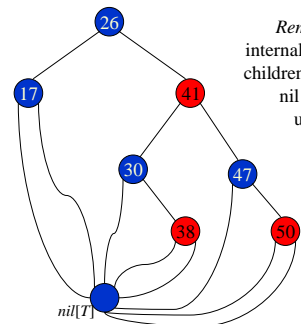
Red-black Tree

- ♦ Binary search tree + 1 bit per node: the attribute **color**, which is either **red** or **black**.
- ♦ All other attributes of BSTs are inherited:
 - » **key**, **left**, **right**, and **p**.
- ♦ All empty trees (leaves) are colored black.
 - » We use a single sentinel, **nil**, for all the leaves of red-black tree T , with $color[nil]$ = black.
 - » The root's parent is also $nil[T]$.

Red-black Properties

1. Every node is either **red** or **black**.
2. The **root** is **black**.
3. Every **leaf (nil)** is **black**.
4. If a node is **red**, then both its children are **black**.
5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes.

Red-black Tree – Example



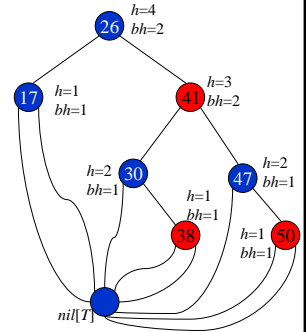
Remember: every internal node has two children, even though nil leaves are not usually shown.

Height of a Red-black Tree

- ♦ **Height of a node:**
 - » $h(x)$ = number of edges in a longest path to a leaf.
- ♦ **Black-height of a node x , $bh(x)$:**
 - » $bh(x)$ = number of black nodes (including $nil[T]$) on the path from x to leaf, not counting x .
- ♦ Black-height of a red-black tree is the black-height of its root.
 - » By Property 5 (For each node, all paths from the node to descendant leaves contain the same number of black nodes),
 - » black height is well defined.

Height of a Red-black Tree

- ♦ Example:
- ♦ **Height of a node:**
 - $h(x)$ = # of edges in a longest path to a leaf.
- ♦ **Black-height of a node**
 - $bh(x)$ = # of black nodes on path from x to leaf, not counting x .
- ♦ **How are they related?**
 - » $bh(x) \leq h(x) \leq 2 bh(x)$

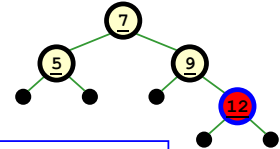


Operations on RB Trees

- ♦ All operations can be performed in $O(\lg n)$ time.
- ♦ The query operations, which don't modify the tree, are performed in exactly the same way as they are in BSTs.
- ♦ Insertion and Deletion are not straightforward.

Red-Black Trees: An Example

- ♦ *Color this tree:*

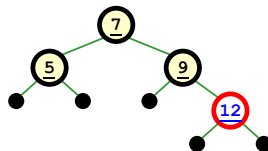


Red-black properties:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendant leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

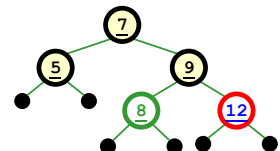
- ♦ Insert 8
 - » *Where does it go?*



1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendant leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

- ♦ Insert 8
 - » *Where does it go?*
 - » *What color should it be?*

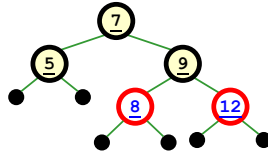


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendant leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 8

- » *Where does it go?*
- » *What color should it be?*

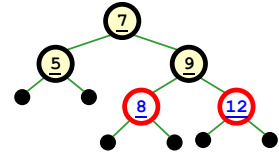


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 11

- » *Where does it go?*

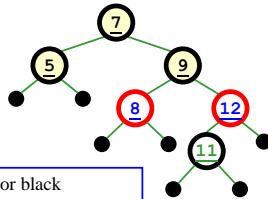


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 11

- » *Where does it go?*
- » *What color?*

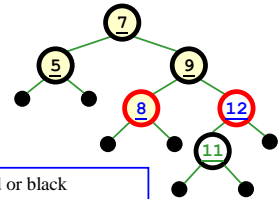


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 11

- » *Where does it go?*
- » *What color?*
 - Can't be red! (#3)

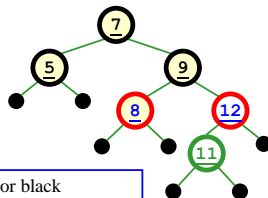


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 11

- » *Where does it go?*
- » *What color?*
 - Can't be red! (#3)
 - Can't be black! (#4)

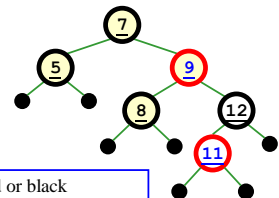


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 11

- » *Where does it go?*
- » *What color?*
 - Solution: recolor the tree

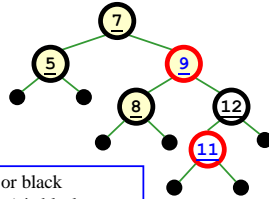


1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 10

» Where does it go?



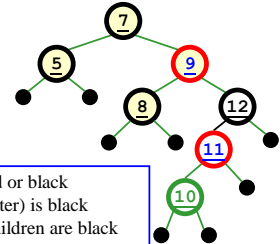
1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

♦ Insert 10

» Where does it go?

» What color?



1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

Red-Black Trees: The Problem With Insertion

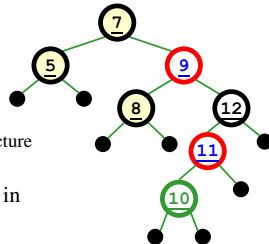
♦ Insert 10

» Where does it go?

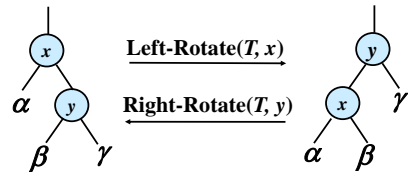
» What color?

- A: no color! Tree is too imbalanced
- Must change tree structure to allow recoloring

» Goal: restructure tree in $O(\lg n)$ time

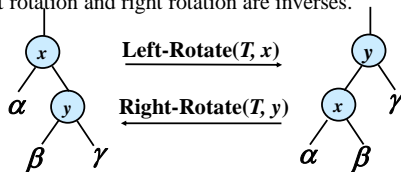


Rotations



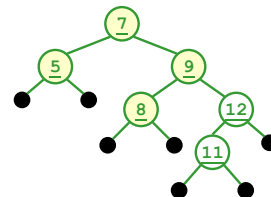
Rotations

- ♦ Rotations are the basic **tree-restructuring** operation for almost all *balanced* search trees.
- ♦ Rotation takes a red-black-tree and a node,
- ♦ Changes pointers to change the local structure, and
- ♦ Won't violate the binary-search-tree property.
- ♦ Left rotation and right rotation are inverses.



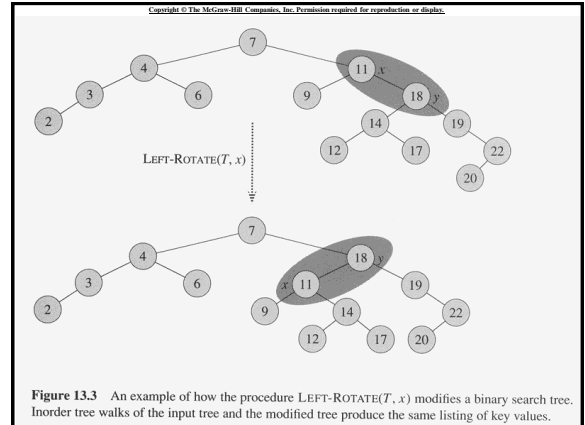
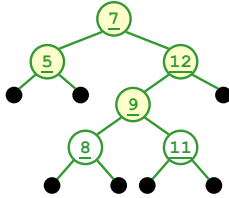
Rotation Example

♦ Rotate left about 9:



Rotation Example

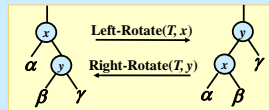
- ♦ Rotate left about 9:



Left Rotation – Pseudo-code

Left-Rotate (T, x)

1. $y \leftarrow \text{right}[x]$ // Set y .
2. $\text{right}[x] \leftarrow \text{left}[y]$ // Turn y 's left subtree into x 's right subtree.
3. **if** $\text{left}[y] \neq \text{nil}[T]$
4. **then** $p[\text{left}[y]] \leftarrow x$
5. $p[y] \leftarrow p[x]$ // Link x 's parent to y .
6. **if** $p[x] = \text{nil}[T]$
7. **then** $\text{root}[T] \leftarrow y$
8. **else if** $x = \text{left}[p[x]]$
9. **then** $\text{left}[p[x]] \leftarrow y$
10. **else** $\text{right}[p[x]] \leftarrow y$
11. $\text{left}[y] \leftarrow x$ // Put x on y 's left.
12. $p[x] \leftarrow y$



Rotation

- ♦ The pseudo-code for Left-Rotate assumes that
 - » $\text{right}[x] \neq \text{nil}[T]$, and
 - » root 's parent is $\text{nil}[T]$.
- ♦ Left Rotation on x , makes x the left child of y , and the left subtree of y into the right subtree of x .
- ♦ Pseudocode for Right-Rotate is symmetric: exchange *left* and *right* everywhere.
- ♦ **Time:** $O(1)$ for both Left-Rotate and Right-Rotate, since a constant number of pointers are modified.

Red-black Properties

1. Every node is either **red** or **black**.
2. The **root** is **black**.
3. Every **leaf** (*nil*) is **black**.
4. If a node is **red**, then both its children are **black**.
5. For each node, all paths from the node to descendant leaves contain the **same number of black nodes**.

Insertion in RB Trees

- ♦ Insertion must preserve all red-black properties.
- ♦ Should an inserted node be colored **Red?** **Black?**
- ♦ **Basic steps:**
 - » Use Tree-Insert from BST (slightly modified) to insert a node x into T .
 - Procedure **RB-Insert(x)**.
 - » **Color** the **node x red**.
 - » Fix the modified tree by **re-coloring nodes** and **performing rotation** to **preserve RB tree property**.
 - Procedure **RB-Insert-Fixup**.

Insertion

RB-Insert(T, z)

```

1.  $y \leftarrow \text{nil}[T]$ 
2.  $x \leftarrow \text{root}[T]$ 
3. while  $x \neq \text{nil}[T]$ 
4.   do  $y \leftarrow x$ 
5.     if  $\text{key}[z] < \text{key}[x]$ 
6.       then  $x \leftarrow \text{left}[x]$ 
7.       else  $x \leftarrow \text{right}[x]$ 
8.  $p[z] \leftarrow y$ 
9. if  $y = \text{nil}[T]$ 
10.  then  $\text{root}[T] \leftarrow z$ 
11.  else if  $\text{key}[z] < \text{key}[y]$ 
12.    then  $\text{left}[y] \leftarrow z$ 
13.    else  $\text{right}[y] \leftarrow z$ 

```

RB-Insert(T, z) Contd.

```

14.  $\text{left}[z] \leftarrow \text{nil}[T]$ 
15.  $\text{right}[z] \leftarrow \text{nil}[T]$ 
16.  $\text{color}[z] \leftarrow \text{RED}$ 
17. RB-Insert-Fixup( $T, z$ )

```

Which of the RB properties might be violated?

Fix the violations by calling RB-Insert-Fixup.

Red-black Properties

1. Every node is either **red** or **black**.
2. The **root** is **black**.
3. Every **leaf** (*nil*) is **black**.
4. If a node is **red**, then both its children are **black**.
5. For each node, all paths from the node to descendant leaves contain the **same number of black nodes**.

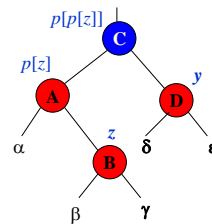
Properties violations

- ♦ **Property 1** (each node black or red): hold
- ♦ **Property 3** (each leaf is black sentinel): hold.
- ♦ **Property 5**: same number of blacks: hold
- ♦ **Property 2**: (root is black), not, if z is root (and colored red).
- ♦ **Property 4**: (the child of a red node must be black), not, if z 's parent is red.

Insertion – Fixup

- ♦ Problem: **We may have one pair of consecutive reds where we did the insertion.**
- ♦ Solution: **Rotate the tree and then Three cases have to be handled**

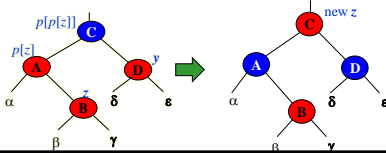
- ♦ **Case 1, 2, 3**: $p[z]$ is the left child of $p[p[z]]$.
- ♦ Correspondingly, **there are 3 other cases**,
- ♦ In which $p[z]$ is the right child of $p[p[z]]$



Insertion – Fixup

RB-Insert-Fixup(T, z)

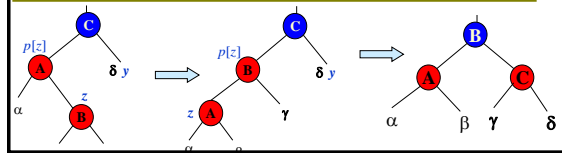
1. **while** $color[p[z]] = \text{RED}$
2. **do if** $p[z] = \text{left}[p[p[z]]]$
3. **then** $y \leftarrow \text{right}[p[p[z]]]$
4. **if** $color[y] = \text{RED}$
5. **then** $color[p[z]] \leftarrow \text{BLACK}$ // Case 1
6. $color[y] \leftarrow \text{BLACK}$ // Case 1
7. $color[p[p[z]]] \leftarrow \text{RED}$ // Case 1
8. $z \leftarrow p[p[z]]$ // Case 1



Insertion – Fixup

RB-Insert-Fixup(T, z) (Contd.)

9. **else if** $z = \text{right}[p[z]]$ // $color[y] \neq \text{RED}$
10. **then** $z \leftarrow p[z]$ // Case 2
11. $\text{LEFT-ROTATE}(T, z)$ // Case 2
12. $color[p[z]] \leftarrow \text{BLACK}$ // Case 3
13. $color[p[p[z]]] \leftarrow \text{RED}$ // Case 3
14. $\text{RIGHT-ROTATE}(T, p[p[z]])$ // Case 3
15. **else** (if $p[z] = \text{right}[p[p[z]]]$) (same as then clause with “right” and “left” exchanged)
- 16.
17. $color[\text{root}[T]] \leftarrow \text{BLACK}$



Insertion – Fixup

Termination

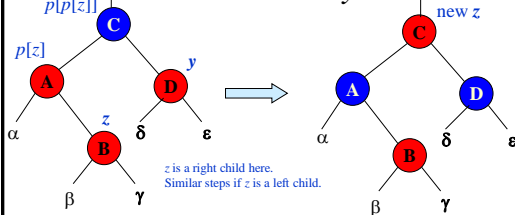
The loop terminates only if $p[z]$ is black. Hence, property 4 is OK.
The last line ensures property 2 always holds.

Copyright © The McGraw-Hill Companies, Inc. Permission is granted for reproduction or display.

RB-INSERT-FIXUP(T, z)

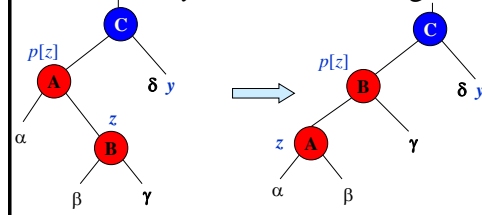
- 1 **while** $color[p[z]] = \text{RED}$
- 2 **do if** $p[z] = \text{left}[p[p[z]]]$
- 3 **then** $y \leftarrow \text{right}[p[p[z]]]$
- 4 **if** $color[y] = \text{RED}$
- 5 **then** $color[p[z]] \leftarrow \text{BLACK}$ ▷ Case 1
- 6 $color[y] \leftarrow \text{BLACK}$ ▷ Case 1
- 7 $color[p[p[z]]] \leftarrow \text{RED}$ ▷ Case 1
- 8 $z \leftarrow p[p[z]]$ ▷ Case 1
- 9 **else if** $z = \text{right}[p[p[z]]]$
- 10 **then** $z \leftarrow p[p[z]]$ ▷ Case 2
- 11 $\text{LEFT-ROTATE}(T, z)$ ▷ Case 2
- 12 $color[p[p[z]]] \leftarrow \text{BLACK}$ ▷ Case 3
- 13 $color[p[p[p[z]]]] \leftarrow \text{RED}$ ▷ Case 3
- 14 $\text{RIGHT-ROTATE}(T, p[p[p[z]]])$ ▷ Case 3
- 15 **else** (same as then clause with “right” and “left” exchanged)
- 16 $color[\text{root}[T]] \leftarrow \text{BLACK}$

Case 1 – uncle y is red



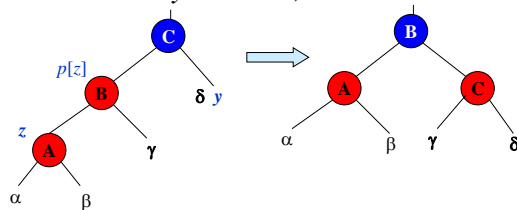
- $p[p[z]]$ (z 's grandparent) must be black, since z and $p[z]$ are both red and there are no other violations of property 4.
- Make $p[z]$ and y black \Rightarrow now z and $p[z]$ are not both red. **But property 5 might now be violated.**
- Make $p[p[z]]$ red \Rightarrow restores property 5.
- The next iteration has $p[p[z]]$ as the new z (i.e., z moves up 2 levels).

Case 2 – y is black, z is a right child



- Left rotate around $p[z]$, $p[z]$ and z switch roles \Rightarrow **now z is a left child, and both z and $p[z]$ are red.**
- Takes us immediately to case 3.

Case 3 – y is black, z is a left child



- ♦ Make $p[z]$ black and $p[p[z]]$ red.
- ♦ Then right rotate on $p[p[z]]$. Ensures **property 4 is maintained**.
- ♦ No longer have 2 reds in a row.
- ♦ $p[z]$ is now black \Rightarrow no more iterations.

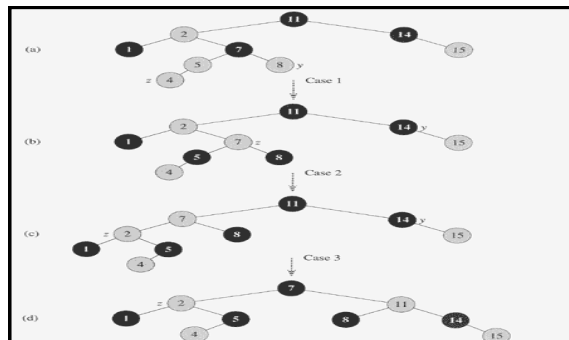


Figure 13.4 The operation of RB-INSERT-FIXUP. (a) A node z after insertion. Since z and its parent $p[z]$ are both red, a violation of property 4 occurs. Since z 's uncle y is red, case 1 in the code can be applied. Nodes are recolored and the pointer z is moved up the tree, resulting in the tree shown in (b). Once again, z and its parent are both red, but z 's uncle y is black. Since z is the right child of $p[z]$, case 2 can be applied. A left rotation is performed, and the tree that results is shown in (c). Now z is the left child of its parent, and case 3 can be applied. A right rotation yields the tree in (d), which is a legal red-black tree.

Algorithm Analysis

- ♦ $O(\lg n)$ time to get through RB-Insert up to the call of RB-Insert-Fixup.
- ♦ Within **RB-Insert-Fixup**:
 - » Each iteration takes $O(1)$ time.
 - » Each iteration but the last **moves z up 2 levels**.
 - » $O(\lg n)$ levels $\Rightarrow O(\lg n)$ time.
 - » Thus, insertion in a red-black tree takes $O(\lg n)$ time.
 - » There are at most 2 rotations overall.

Deletion

- ♦ Deletion, like insertion, should preserve all the RB properties.
- ♦ The properties that may be violated depends on the color of the deleted node.
 - » **Red – OK. Why?**
 - » **Black?**
- ♦ Steps:
 - » Do regular BST deletion.
 - » **Fix any violations of RB properties** that may result.

Deletion

RB-Delete(T, z)

1. if $\text{left}[z] = \text{nil}[T]$ or $\text{right}[z] = \text{nil}[T]$
2. then $y \leftarrow z$
3. else $y \leftarrow \text{TREE-SUCCESSOR}(z)$
4. if $\text{left}[y] = \text{nil}[T]$
5. then $x \leftarrow \text{left}[y]$
6. else $x \leftarrow \text{right}[y]$
7. $p[x] \leftarrow p[y]$ // Do this, even if x is $\text{nil}[T]$

Deletion

RB-Delete(T, z) (Contd.)

8. if $p[y] = \text{nil}[T]$
9. then $\text{root}[T] \leftarrow x$
10. else if $y = \text{left}[p[y]]$
11. then $\text{left}[p[y]] \leftarrow x$
12. else $\text{right}[p[y]] \leftarrow x$
13. if $y = z$
14. then $\text{key}[z] \leftarrow \text{key}[y]$
15. copy y 's satellite data into z
16. if $\text{color}[y] = \text{BLACK}$
17. then RB-Delete-Fixup(T, x)
18. return y

The node passed to the fixup routine is the lone child of the spliced up node, or the sentinel.

RB Properties Violation

- ♦ If y is black, we could have violations of red-black properties:
 - » Prop. 1. OK.
 - » Prop. 2. If y is the root and x is red, then the root has become red.
 - » Prop. 3. OK.
 - » Prop. 4. Violation if $p[y]$ and x are both red.
 - » Prop. 5. Any path containing y now has 1 fewer black node.

RB Properties Violation

- ♦ Prop. 5. Any path containing y now has 1 fewer black node.
 - » Correct by giving x an “extra black.”
 - » Add 1 to count of black nodes on paths containing x .
 - » Now property 5 is OK, but property 1 is not.
 - » x is either *doubly black* (if $color[x] = \text{BLACK}$) or *red & black* (if $color[x] = \text{RED}$).
 - » The attribute $color[x]$ is still either RED or BLACK. No new values for $color$ attribute.
 - » In other words, the extra blackness on a node is by virtue of x pointing to the node.
- ♦ Remove the violations by calling RB-Delete-Fixup.

Deletion – Fixup

RB-Delete-Fixup(T, x)

1. while $x \neq \text{root}[T]$ and $color[x] = \text{BLACK}$
2. do if $x = \text{left}[p[x]]$
3. then $w \leftarrow \text{right}[p[x]]$
4. if $color[w] = \text{RED}$
5. then $color[w] \leftarrow \text{BLACK}$ // Case 1
6. $color[p[x]] \leftarrow \text{RED}$ // Case 1
7. LEFT-ROTATE($T, p[x]$) // Case 1
8. $w \leftarrow \text{right}[p[x]]$ // Case 1

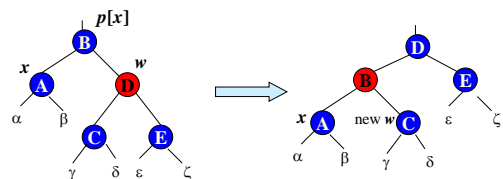
RB-Delete-Fixup(T, x) (Contd.)

- /* x is still $\text{left}[p[x]]$ */
9. if $color[\text{left}[w]] = \text{BLACK}$ and $color[\text{right}[w]] = \text{BLACK}$
 10. then $color[w] \leftarrow \text{RED}$ // Case 2
 11. $x \leftarrow p[x]$ // Case 2
 12. else if $color[\text{right}[w]] = \text{BLACK}$
 13. then $color[\text{left}[w]] \leftarrow \text{BLACK}$ // Case 3
 14. $color[w] \leftarrow \text{RED}$ // Case 3
 15. RIGHT-ROTATE(T, w) // Case 3
 16. $w \leftarrow \text{right}[p[x]]$ // Case 3
 17. $color[w] \leftarrow color[p[x]]$ // Case 4
 18. $color[p[x]] \leftarrow \text{BLACK}$ // Case 4
 19. $color[\text{right}[w]] \leftarrow \text{BLACK}$ // Case 4
 20. LEFT-ROTATE($T, p[x]$) // Case 4
 21. $x \leftarrow \text{root}[T]$ // Case 4
 22. else (same as then clause with “right” and “left” exchanged)
 23. $color[x] \leftarrow \text{BLACK}$

Deletion – Fixup

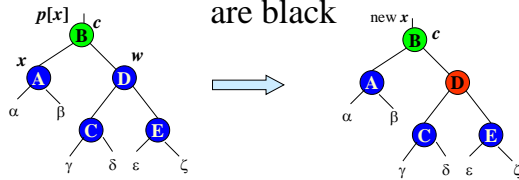
- ♦ **Idea:** Move the extra black up the tree until x points to a red & black node \Rightarrow turn it into a black node,
- ♦ x points to the root \Rightarrow just remove the extra black, or
- ♦ We can do certain rotations and recolorings and finish.
- ♦ Within the **while** loop:
 - » x always points to a nonroot doubly black node.
 - » w is x 's sibling.
 - » w cannot be $\text{nil}[T]$, since that would violate property 5 at $p[x]$.
- ♦ 8 cases in all, 4 of which are symmetric to the other.

Case 1 – w is red



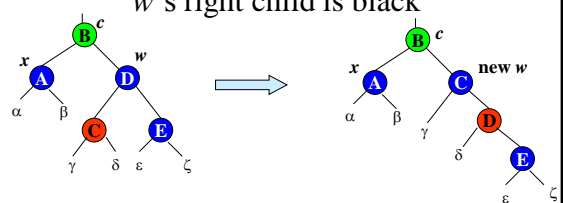
- ♦ w must have black children.
- ♦ Make w black and $p[x]$ red (because w is red $p[x]$ couldn't have been red).
- ♦ Then left rotate on $p[x]$.
- ♦ New sibling of x was a child of w before rotation \Rightarrow must be black.
- ♦ Go immediately to case 2, 3, or 4.

Case 2 – w is black, both w 's children are black



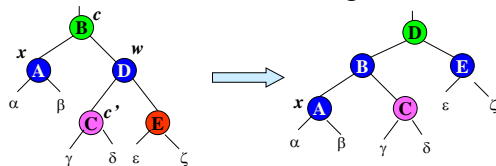
- ♦ Take 1 black off x (\Rightarrow singly black) and off w (\Rightarrow red).
- ♦ Move that black to $p[x]$.
- ♦ Do the next iteration with $p[x]$ as the new x .
- ♦ If entered this case from case 1, then $p[x]$ was red \Rightarrow new x is red & black \Rightarrow color attribute of new x is RED \Rightarrow loop terminates. Then new x is made black in the last line.

Case 3 – w is black, w 's left child is red, w 's right child is black



- ♦ Make w red and w 's left child black.
- ♦ Then right rotate on w .
- ♦ New sibling w of x is black with a red right child \Rightarrow case 4.

Case 4 – w is black, w 's right child is red



- ♦ Make w be $p[x]$'s color (c).
- ♦ Make $p[x]$ black and w 's right child black.
- ♦ Then left rotate on $p[x]$.
- ♦ Remove the extra black on x (\Rightarrow x is now singly black) without violating any red-black properties.
- ♦ All done. Setting x to root causes the loop to terminate.

Analysis

- ♦ $O(\lg n)$ time to get through RB-Delete up to the call of RB-Delete-Fixup.
- ♦ Within RB-Delete-Fixup:
 - » Case 2 is the only case in which more iterations occur.
 - x moves up 1 level.
 - Hence, $O(\lg n)$ iterations.
 - » Each of cases 1, 3, and 4 has 1 rotation $\Rightarrow \leq 3$ rotations in all.
 - » Hence, $O(\lg n)$ time.